

# Decentralized computing and personal data stores

**Nicolas Seydoux, PhD**

Developer at Inrupt, Solid community member

# Introductory use-case: Reasoning in "smart" homes

# Typical smart home use case

- Remote control for connected devices
- Data gathering, **analytics and reporting**
- **Scenario-based** automation

# Use case: smart home deployment

- The user has deployed:
  - Smart plugs, to monitor power consumption
  - Smoke detectors
  - Connected thermostat and weather station
  - Lighting system with presence sensors...


# Use-case apps

- App A: Rule-based reasoning
  - Light schedule scenario
  - Smoke detection
  - Temperature preferences
- App B: Machine-learning based
  - Power consumption profile

# Typical approach and its limitations

# Cloud-based IoT service platforms

- Data is sent from local network to remote platform
- Cloud-based platform allows automation, analytics, remote control...



# Service provider dependency

- The data is controlled by the service provider
- The user cannot migrate easily to another platform





# Privacy concerns

- Smart home data is collected in an intimate setting.
- Potential inference of personal information:
  - Presence patterns at home
  - Religious practices

# Interoperability issues


- The data is controlled by the app
- Systems may be isolated in silos

# Security vulnerabilities

- Confidentiality
  - Concentration of many users data in one place
  - Data must be sent to a third party
- Availability
  - Single point of failure

# Shifting from Cloud to Fog computing

- The user keeps the data on premise
- Dispatching computing instead of data
- Inference results flow in the hierarchy too



# Shifting from app-centric to user-centric

- Each user is in control of their data
- The data store may be self-hosted, or hosted by a trusted party
- Migrating from one store to another is possible
- Applications read data from the user-controlled store

# Solid, access control on top of Linked Data

# The user

- They control the data.
- They use apps to get services.



# The app

- It consumes the data: **reasoning**.
- It produces more data: **inference**.





# The Pod

- It hosts the data.
- It enforces access permissions.




# The Identity Provider

- It authenticates the user.



# The WebID

- An IRI identifying the user.
- WebID profile links to user data.



# Authentication patterns


# Two patterns to access data

- The User logs in the Client
- The Client acts on its own behalf (bot)

# Authentication patterns: User login

# User initiates login

- The Client doesn't manage the User's identity
- The Client dedirects the User to the IdP



# User logs in

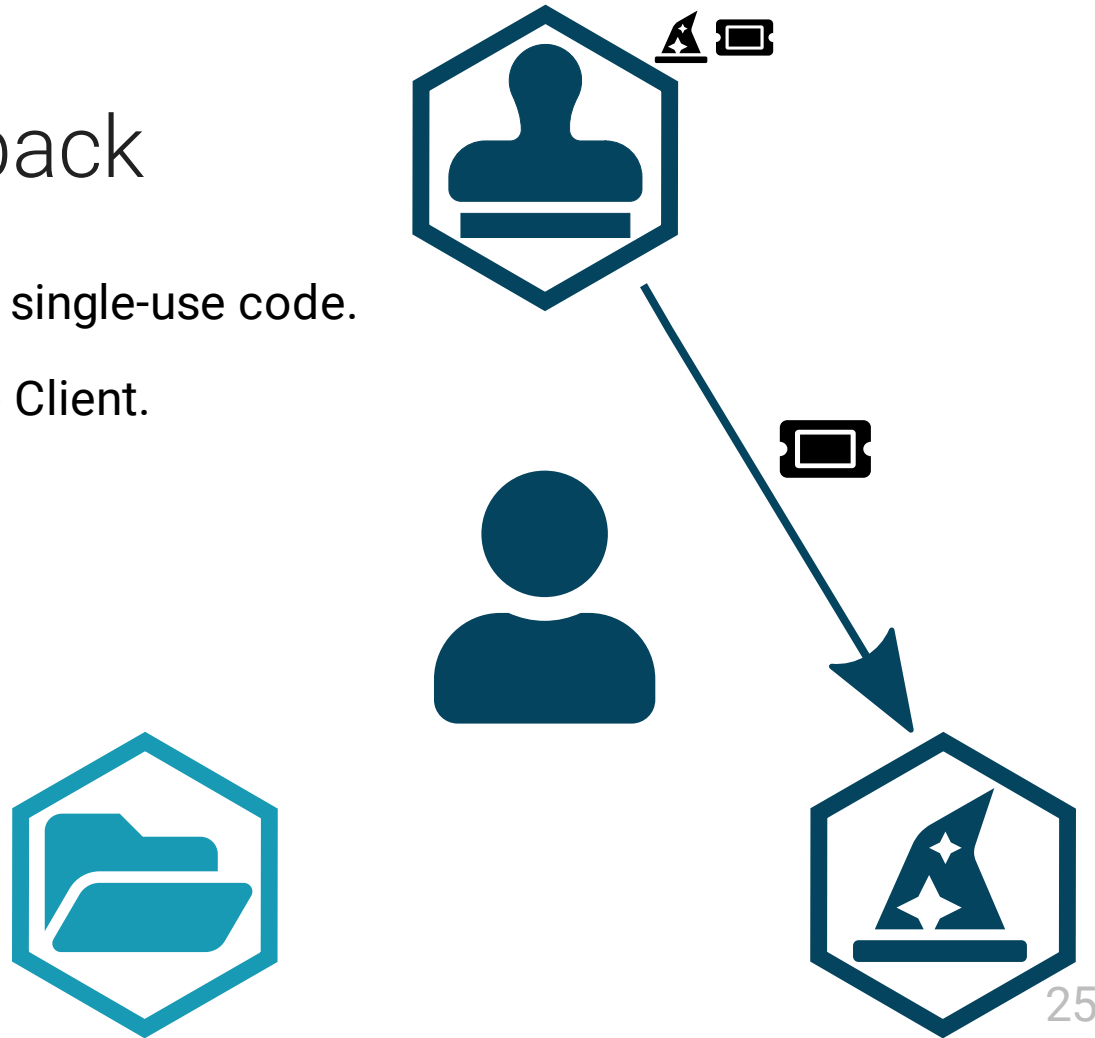
- The authentication method is out of scope.
- The Client never sees the User's credentials.






# IdP redirects User back

- The IdP provides the Client with a single-use code.
- The IdP associates the code to the Client.




# Client sends code

- The Client sends the code to the IdP




# IdP sends tokens

- The ID token is meant for the Client.
- The Access Token is meant for the Pod.



# Authenticated access


- The Client acts on behalf of the User.
- The Client has its own identity too.



# Authentication patterns: Client login


# Developer registers Client

- The IdP has information about the Client
- The Client gets credentials.




# Client logs in

- The Client exchanges its credentials for tokens to the IdP



# Authenticated access


- The Client acts on its own behalf.





# Navigating the data

# Discovering a user's data store



# Getting the Pod root

# Traversing to a container

# Getting an RDF resource

# Discovering more resources


# Getting a non-RDF resource

# Reasoning against Solid Pods



# Centralizing common, decentralizing individual

- Not all data is personal
  - Vocabularies/instances
  - Devices datasheets/local deployments
  - "Common sense"/user preferences
- **Private data** linked to **shared knowledge bases**



# Reasoning patterns: client side

- Since the data is directly available to the client, all of the reasoning may happen in the browser
- Lightweight for the service provider
- Entirely user-centric
- User login authentication pattern

# Client-side reasoning use case

- Suitable for scenarios centered on one user data
- Power consumption analytics, complex event processing...
- The user is logged in the app doing the reasoning

# Reasoning patterns: server-side

- Data from multiple users is aggregated in a single graph
- Allows to identify patterns beyond individual users
- Bot login authentication pattern


# Server-side reasoning use case

- Suitable for scenarios aggregating data from multiple users
- Machine learning training, large-scale predictions...
- Individual user data is collected by the app doing the reasoning

Taking advantage of  
app/data separation

# Beyond app silos

- All data for a user is available in their Pod
- Not all apps need to harvest: data is collected once, reused infinitely
- Power consumption data may be linked to thermostat and weather



# Decentralizing chained inferences

- A Pod is a read/write storage
- Inferences on user data should be written back to the user's Pod
- Inferences chained through the Pod instead of sequentially



# Challenges and discussion topics

# Interoperability, symbolic and numeric AI

- Symbolic reasoning is a natural fit for Linked Data
- Numerical approaches extremely powerful too, but act as black boxes.
- How to make numerical models interoperable?

# Data discoverability

- Data stored by Pod providers
- A given app or user may not have access to some data from a source
- User-centric apps only need one user's data
- How to find many small datasets to train a model?

# From individual to collective and back

- User data needs to be consolidated into a single dataset for ML training
- Performance issue:  $N$  data stores  $\Rightarrow N$  requests
- Is it possible to adapt the result back to each individual user?

# Maintaining provenance

- A model is trained based on data from dynamic sources
- What provenance information should be captured to keep the model accurate?

# Preserving privacy in sharing

- Anonymizing data before sharing it
- Temporal or spatial aggregation at different scales

# Conclusion

# Reasoning and personal data stores

- Solid is a set of specifications centering the data around the user
- Different decentralized processing approaches are complementary
- Decoupling data from applications offers new opportunities





# Bibliography

- A. Coburn, E. Pavlik, et D. Zagidulin, « Solid-OIDC ». [Online] 2022. <https://solidproject.org/TR/oidc>
- S. Capadisli, T. Berners-Lee, R. Verborgh, and K. Kjernsmo, “Solid Protocol,” 2021. <https://solidproject.org/TR/protocol>.
- N. Seydoux, K. Drira, N. Hernandez, et T. Monteil, « EDR: A Generic Approach for the Distribution of Rule-Based Reasoning in a Cloud-Fog continuum », Semantic Web Journal, 2019.
- M. Terdjimi, L. Médini, et M. Mrissa, « Web Reasoning Using Fact Tagging », in Companion Proceedings of the The Web Conference 2018.

- Credit to fontawesome for the icons ([license](#))